

Object Detection methods for wildlife species monitoring

Yulia Pavlova

project report

9 September 2019

The objective of the task was to create a program to classify wild-life animals from the photos taken by hunter cameras. The program (the first version of it) is using a TensorFlow-trained classifier to detection. The model was retrained for the specified task using images from Loimaa 2017, plus additional data of species called 'Roe deer'. I have used Single Shot Multibox Detector (SSD) neural network system.

The results of object detector model are shown below (Fig.1-Fig.4). The program draws boxes around detected animals and identifies species of the animals.

Fig. 1.



Fig.2




Fig.3


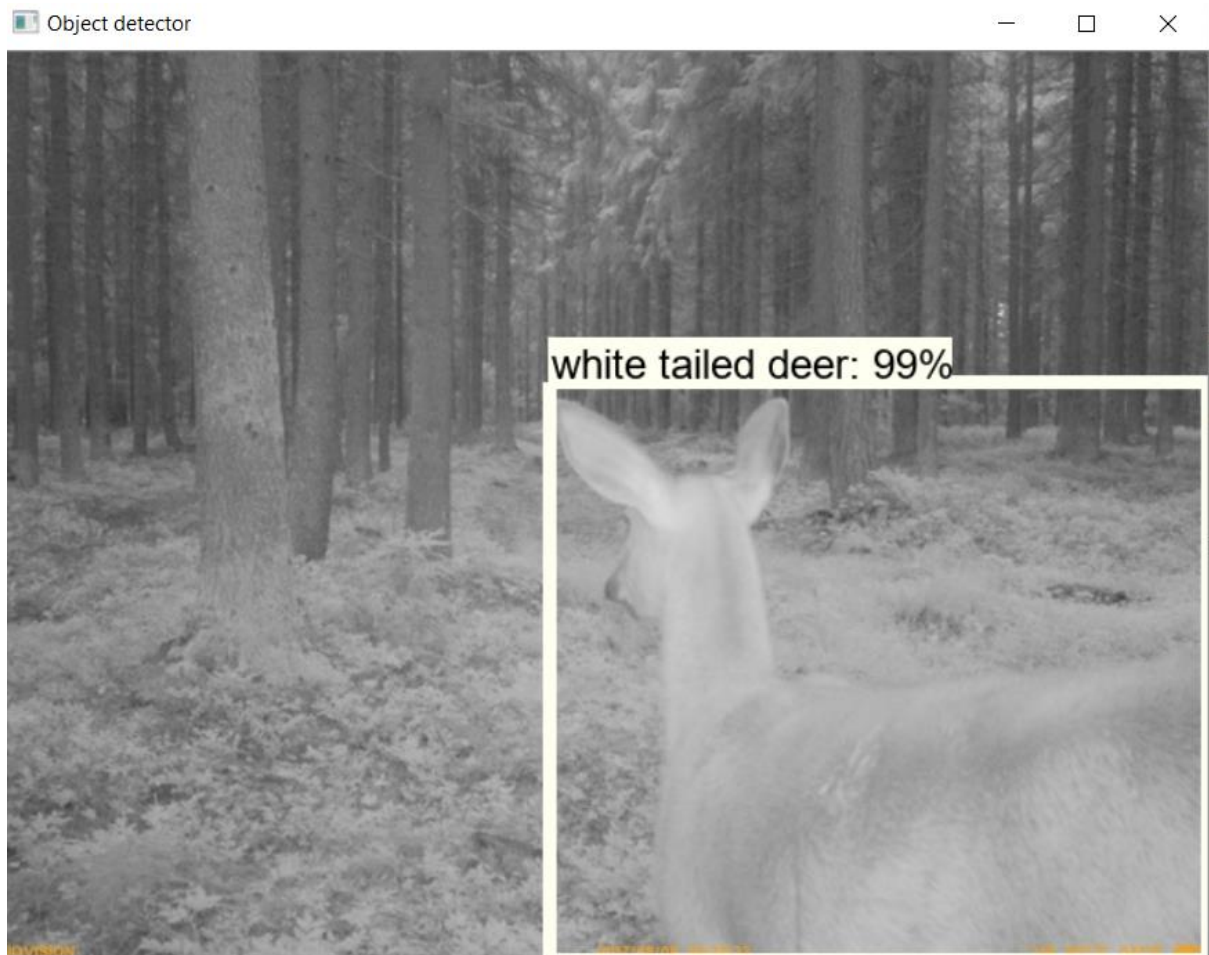


Fig.4



Initial observations show that the part model of identification the species works well, however, detecting part of model that draws boxes has some room for improvements: in many cases it fails to detect animals.

Secondly, the work on describing statistics of the model is not completed. This is some work for the future.

Thirdly, the detection can be performed on the video stream (the code is also provided).

Below I describe the steps and provide the useful links.

Installation

<https://tensorflow-object-detection-api-tutorial.readthedocs.io/en/latest/install.html>

An important start of the project was related to installation of TensorFlow object detection API. The link above provides all necessary steps of the process. Administrative rights are necessary. An optional but highly recommended prerequisite is Anaconda tool that helps to manage virtual environments. For this task I have chosen GPU version of TensorFlow, which is commonly considered to be faster. However, installation process is somewhat more demanding due to need of additional installation of Graphics and CUDE drivers.

When installation process is all set up, it is advisable to create a new directory – ‘wild-life_animal_detection’, which contains annotations, images (test/train), pre-trained-model, training, prog-utils, readme.md (contains project description). This is done for organizational purposes.

- data
- images (test/train)
- object_detection
 - o training
 - o inference_graph
- readme.md (contains project description).

Gather data, data preprocessing and labelling steps

Now that the TensorFlow Object Detection API is all set up and ready to go, we need to provide the images it will use to train a new detection classifier. TensorFlow needs hundreds of images of an object to train a good detection classifier. To train a robust classifier, the training images should have random objects in the image along with the desired objects and should have a variety of backgrounds and lighting conditions. There should be some images where the desired object is partially obscured, overlapped with something else, or only halfway in the picture.

Make sure the images aren't too large. They should be less than 200KB each, and their resolution shouldn't be more than 720x1280. The larger the images are, the longer it will take to train the classifier. *Resizer.py* script (in prog_utils folder) has been used to reduce the size of the images.

Annotations

With all the pictures gathered, it's time to label the desired objects in every picture. *Labellmg* is a great tool for labeling images, and its GitHub page has very clear instructions on how to install and use it . Download and install *Labellmg*

<https://tensorflow-object-detection-api-tutorial.readthedocs.io/en/latest/install.html#labelimg-install>

Open each folder, select a photo, draw a box around each animal in each image and assign a tag. Repeat the process for all the images in the directory. Labelimg saves a .xml file containing the label data for each image. These .xml files then later are used to generate TFRecords, which are one of the inputs to the TensorFlow trainer. Once all images are labeled and saved each image, there will be one .xml file for each image.

After that, I randomly move 10% of them (images and .xml files) to the images\test directory, and 90% of them to the images\train directory. For that I use another script *move_split.py* (in prog_utils folder).

Training

With the images labeled, need to generate the TFRecords that serve as input data to the TensorFlow training model. Here I used the *xml_to_csv.py* and *generate_tfrecord.py* scripts, with some slight modifications, from

https://github.com/datitran/raccoon_dataset

First, the image .xml data will be used to create .csv files containing all the data for the train and test images. From the \object_detection folder, issue the following command in the Anaconda command prompt: `python xml_to_csv.py`.

This creates a *train_labels.csv* and *test_labels.csv* file in the \object_detection\data folder.

Next, open the *generate_tfrecord.py* file in a text editor. Replace the label map starting at with your own label map, where each object is assigned an ID number. This same number assignment will be used when configuring the *animal_mask.pbtxt*. In this exercise, I use four tags: rabbit, roe, moose, and white tailed deer (there is not enough images of other species from Loimaa 2017).

Then, generate the TFRecord files by issuing these commands from the \object_detection folder:

```
python generate_tfrecord.py --csv_input=images\train_labels.csv --
image_dir=images\train --output_path=train.record
python generate_tfrecord.py --csv_input=images\test_labels.csv --
image_dir=images\test --output_path=test.record
```

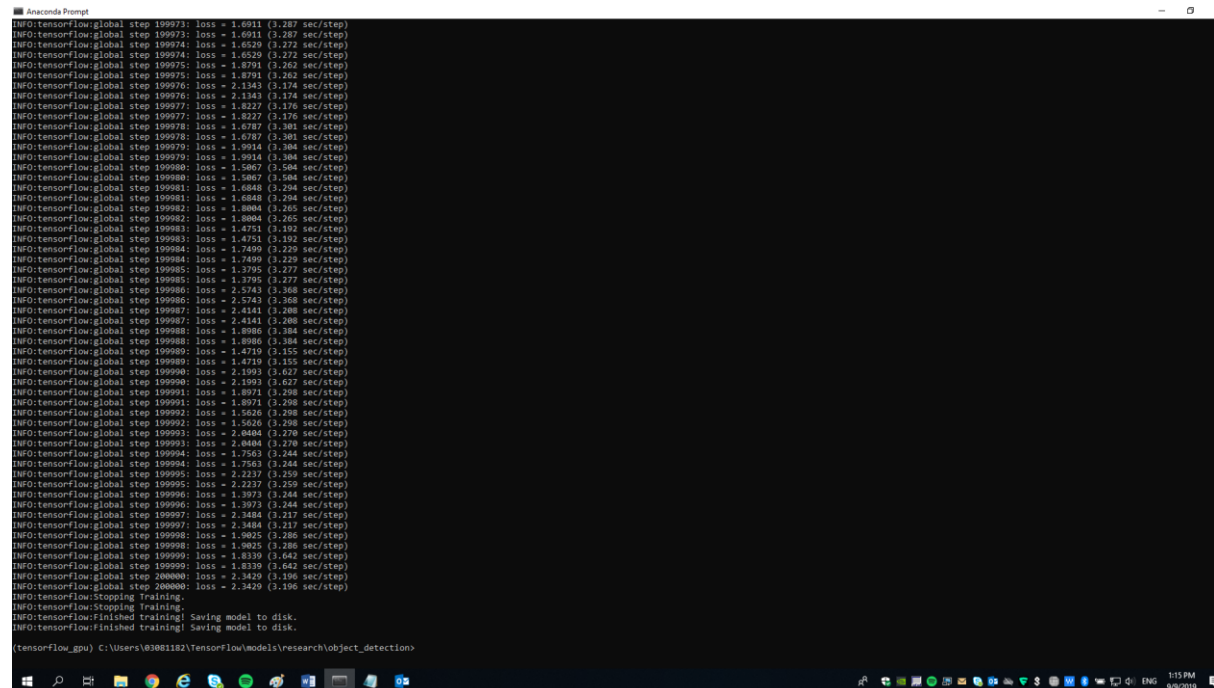
https://github.com/tensorflow/models/blob/master/research/object_detection/object_detection_tutorial.ipynb

Finally, the object detection training pipeline is configured. It defines which model and what parameters are used for training. Navigate to \object_detection\samples\configs and copied the *ssd_inception_v2_pets.config* file into the \object_detection\training directory.

After that training of the model can begin using *train.py* (in anaconda prompt from object_detection folder)

```
python train.py --logtostderr --train_dir=training/ --  
pipeline_config_path=training/faster_rcnn_inception_v2_pets.config
```

Training takes several days (about 5 days), I ran it up to 200,000 steps.



```
INFO:tensorflow:global step 199973: loss = 1.6011 (3.287 sec/step)  
INFO:tensorflow:global step 199974: loss = 1.6011 (3.287 sec/step)  
INFO:tensorflow:global step 199975: loss = 1.6029 (3.272 sec/step)  
INFO:tensorflow:global step 199976: loss = 1.6529 (3.272 sec/step)  
INFO:tensorflow:global step 199977: loss = 1.8791 (3.262 sec/step)  
INFO:tensorflow:global step 199978: loss = 1.8791 (3.262 sec/step)  
INFO:tensorflow:global step 199979: loss = 2.1343 (3.174 sec/step)  
INFO:tensorflow:global step 199980: loss = 2.1343 (3.174 sec/step)  
INFO:tensorflow:global step 199981: loss = 1.8227 (3.176 sec/step)  
INFO:tensorflow:global step 199982: loss = 1.8227 (3.176 sec/step)  
INFO:tensorflow:global step 199983: loss = 1.6787 (3.301 sec/step)  
INFO:tensorflow:global step 199984: loss = 1.6787 (3.301 sec/step)  
INFO:tensorflow:global step 199985: loss = 1.9914 (3.304 sec/step)  
INFO:tensorflow:global step 199986: loss = 1.9914 (3.304 sec/step)  
INFO:tensorflow:global step 199987: loss = 1.5867 (3.504 sec/step)  
INFO:tensorflow:global step 199988: loss = 1.5867 (3.504 sec/step)  
INFO:tensorflow:global step 199989: loss = 1.6848 (3.294 sec/step)  
INFO:tensorflow:global step 199990: loss = 1.6848 (3.294 sec/step)  
INFO:tensorflow:global step 199991: loss = 1.8804 (3.265 sec/step)  
INFO:tensorflow:global step 199992: loss = 1.8804 (3.265 sec/step)  
INFO:tensorflow:global step 199993: loss = 1.4751 (3.192 sec/step)  
INFO:tensorflow:global step 199994: loss = 1.4751 (3.192 sec/step)  
INFO:tensorflow:global step 199995: loss = 1.7499 (3.220 sec/step)  
INFO:tensorflow:global step 199996: loss = 1.7499 (3.220 sec/step)  
INFO:tensorflow:global step 199997: loss = 1.7899 (3.229 sec/step)  
INFO:tensorflow:global step 199998: loss = 1.7899 (3.229 sec/step)  
INFO:tensorflow:global step 199999: loss = 1.3795 (3.277 sec/step)  
INFO:tensorflow:global step 200000: loss = 1.3795 (3.277 sec/step)  
INFO:tensorflow:global step 200001: loss = 2.5743 (3.368 sec/step)  
INFO:tensorflow:global step 200002: loss = 2.5743 (3.368 sec/step)  
INFO:tensorflow:global step 200003: loss = 2.4141 (3.208 sec/step)  
INFO:tensorflow:global step 200004: loss = 2.4141 (3.208 sec/step)  
INFO:tensorflow:global step 200005: loss = 1.8986 (3.384 sec/step)  
INFO:tensorflow:global step 200006: loss = 1.8986 (3.384 sec/step)  
INFO:tensorflow:global step 200007: loss = 1.4719 (3.155 sec/step)  
INFO:tensorflow:global step 200008: loss = 1.4719 (3.155 sec/step)  
INFO:tensorflow:global step 200009: loss = 2.1993 (3.627 sec/step)  
INFO:tensorflow:global step 200010: loss = 2.1993 (3.627 sec/step)  
INFO:tensorflow:global step 200011: loss = 1.8971 (3.298 sec/step)  
INFO:tensorflow:global step 200012: loss = 1.8971 (3.298 sec/step)  
INFO:tensorflow:global step 200013: loss = 1.5626 (3.298 sec/step)  
INFO:tensorflow:global step 200014: loss = 1.5626 (3.298 sec/step)  
INFO:tensorflow:global step 200015: loss = 2.0444 (3.270 sec/step)  
INFO:tensorflow:global step 200016: loss = 2.0444 (3.270 sec/step)  
INFO:tensorflow:global step 200017: loss = 1.7563 (3.244 sec/step)  
INFO:tensorflow:global step 200018: loss = 1.7563 (3.244 sec/step)  
INFO:tensorflow:global step 200019: loss = 1.2237 (3.259 sec/step)  
INFO:tensorflow:global step 200020: loss = 1.2237 (3.259 sec/step)  
INFO:tensorflow:global step 200021: loss = 1.3973 (3.244 sec/step)  
INFO:tensorflow:global step 200022: loss = 1.3973 (3.244 sec/step)  
INFO:tensorflow:global step 200023: loss = 2.3484 (3.217 sec/step)  
INFO:tensorflow:global step 200024: loss = 2.3484 (3.217 sec/step)  
INFO:tensorflow:global step 200025: loss = 1.9825 (3.286 sec/step)  
INFO:tensorflow:global step 200026: loss = 1.9825 (3.286 sec/step)  
INFO:tensorflow:global step 200027: loss = 1.8339 (3.642 sec/step)  
INFO:tensorflow:global step 200028: loss = 1.8339 (3.642 sec/step)  
INFO:tensorflow:global step 200029: loss = 2.3429 (3.196 sec/step)  
INFO:tensorflow:global step 200030: loss = 2.3429 (3.196 sec/step)  
INFO:tensorflow:Stopping Training.  
INFO:tensorflow:Finished training! Saving model to disk.  
INFO:tensorflow:Finished training! Saving model to disk.  
(tensorflow_gpu) C:\Users\W3881182\TensorFlow\models\research\object_detection>
```

Each step of training reports the loss. It will start high and get lower and lower as training progresses. The loss numbers will be different if a different model is used. MobileNet-SSD starts with a loss of about 20, and should be trained until the loss is consistently under 2. We go loss around 2, so some improvement to the dataset is necessary, and also fine-tuning of the model. More recommendations can drawn upon analyzing model's statics.

The progress of the training job can also be viewed by using TensorBoard. To do this, open a new instance of Anaconda Prompt, activate the tensorflow virtual environment, and from object_detection directory, and issue the following command:

```
tensorboard --logdir=training
```

This will create a webpage on your local machine at YourPCName:6006, which can be viewed through a web browser. The TensorBoard page provides information and graphs that show how the training is progressing. One important graph is the Loss graph, which shows the overall loss of the classifier over time (see it in attachment).

TensorBoard

In attachment you can see statistics from the tensorboard. It shows that model converges nicely, even without tuning and altering hyperparameters

https://www.tensorflow.org/guide/summaries_and_tensorboard

Testing the model

Now that training is complete, the last step is to generate the frozen inference graph (.pb file). From the \object_detection folder, issue the following command:

```
python export_inference_graph.py --input_type image_tensor --pipeline_config_path
training/ssd_inception_v2_pets.config --trained_checkpoint_prefix training/model.ckpt-
200000 --output_directory inference_graph
```

This creates a frozen_inference_graph.pb file in the \object_detection\inference_graph folder. The .pb file contains the object detection classifier.

To test your object detector, move a picture of the object or objects into the \object_detection folder, and change the IMAGE_NAME variable in the *Object_detection_image.py* to match the file name of the picture. Alternatively, you can use a video of the objects (using *Object_detection_video.py*), or just plug in a USB webcam and point it at the objects (using *Object_detection_webcam.py*).

To run any of the scripts, type “idle” in the Anaconda Command Prompt (with the “tensorflow” virtual environment activated) and press ENTER. This will open IDLE, and from there, you can open any of the scripts and run them. The object detector will initialize for about 10 seconds and then display a window showing any objects it’s detected in the image. Results of the object_detector_image.py are presented in the beginning of this document Fig. [1—4].

Other information about modern neural networks

Modern methods in object detection

Object detection is a common term for computer vision techniques classifying and locating objects in an image. Modern object detection is largely based on use of convolutional neural networks. Some of the most relevant system types today are Faster R-CNN, R-FCN, Multibox Single Shot Detector (SSD) and YOLO (You Only Look Once).

Original **R-CNN** method worked by running a neural net classifier on samples cropped from images using externally computed box proposals (=samples cropped with externally computed box proposals; feature extraction done on all the cropped samples). This approach was computationally expensive due to many crops; Fast RCNN reduced the computation by doing the feature extraction only once to the whole image and using cropping on the lower layer (=feature extraction only once on the whole image; samples cropped with externally computed box proposals). Faster RCNN goes a step further and used the extracted features to create class-agnostic boxproposals (=feature extraction only once on the whole image; no externally computed box proposals). R-FCN is like Faster R-CNN but the feature cropping is done in a different layer for increased efficiency.

YOLO (You Only Look Once) works on different principle than the aforementioned models: it runs a single convolutional network on the whole input image (once) to predict bounding boxes with confidence scores for each class simultaneously. The advantage of the simplicity of the approach is that the YOLO model is fast (compared to Faster R-CNN and SSD) and it learns a general representation of the objects. This increases localization error rate (also, YOLO does poorly with images with new aspect ratios or small object flocked together) but reduces false positive rate.

Single Shot Multibox Detector (SSD) differs from the R-CNN based approaches by not requiring a second stage per-proposal classification operation. This makes it fast enough for real-time detection applications. However, this comes with a price of reduced precision. Note that “**SSD with MobileNet**” refers to a model where model meta architecture is SSD and the feature extractor type is MobileNet.

Speed-accuracy tradeoff

Many modern object detection applications require real-time speed. Methods such as YOLO or SSD work that fast, but this tends to come with a decrease in accuracy of predictions, whereas models such as Faster R-CNN achieve high accuracy but are more expensive to run. The cost in model speed depends on the application: With larger images (e.g. 600x600) SSD works comparable to more computationally expensive models such as Faster R-CNN, even as on smaller images its performance is considerably lower.

TensorFlow Object Detection API with SSD model with Mobilenet

TensorFlow Object Detection API is “an open source framework built on top of TensorFlow” that aims to make it easy to “construct, train and deploy object detection models”. To achieve this, TensorFlow Object Detection API provides the user with multiple pre-trained object detection models with instructions and example codes for fine-tuning and using the models for object detection tasks. TensorFlow Object Detection API can be used with different pre-trained models. In this work, two models were chosen and tested. A SSD model with Mobilenet (ssd_mobilenet_v1_coco) was the second choice. The original model had been trained using MSCOCO Dataset which consists of 2.5 million labeled instances in 328 000 images, containing 91 object types such as “person”, “cat” or “dog”. The ssd_mobilenet_v1_coco-model is reported to have mean Average Precision (mAP) of 21 on COCO dataset. In this work, the SSD model was tested both as pre-trained without fine tuning and as fine-tuned for our own dataset.¹

☐ Show data download links

☐ Ignore outliers in chart scaling

Tooltip sorting method: default

Smoothing 0.6

Horizontal Axis
STEP RELATIVE

WALL

Runs


Write a regex to filter runs

□ ○ .

Filter tags (regular expressions supported)


LearningRate	1
Loss	2
Losses	5

Loss/classification_loss
tag: Losses/Loss/classification_loss



□ ≡ ⊞ ↓ run to download CSV JSON

Loss/localization_loss
tag: Losses/Loss/localization_loss



□ ≡ ⊞ ↓ run to download CSV JSON

TotalLoss
tag: Losses/TotalLoss

